

The Prisoner’s Sonata:

Modeling Musical Improvisation with Game Theory

Caroline Marcks, Andrew Mendelsohn, and Jayme Woogerd

1 Introduction

In *A Model of Performance, Interaction, and Improvisation*, Paul Hudak outlines a formal model of musical performance, interaction and improvisation based on the idea that a full musical performance can be understood as a set of complex, interrelated interactions.

To motivate this model, he observed that in any performance there exists an interaction between a player and himself: throughout the performance, a good musician will continuously make adjustments to his playing based on what he is hearing from his own instrument. Likewise, in an ensemble or orchestral performance, a player is necessarily affected by the performance of each of the other players in the ensemble. Finally, each player is working off his interpretation of the musical score and, in the case where the players are improvising, the produced performance may actually deviate wildly from what is given in the score.

In his paper, these interactions are termed *mutually recursive processes* where “the recursion captures feedback [and] mutual recursion captures the interaction between players”. In concrete terms, we define these relationships of interaction algebraically:

$$\mathbf{r} = \text{instr}(\text{player } \mathbf{s} \ \mathbf{r})$$

where \mathbf{r} is the realization, the music actually produced by a player, and \mathbf{s} is the score. Hudak then notes that in the context of improvisation, “although the goal of those involved... is generally one of cooperation, there is also a certain amount of conflict.” Given this natural relationship of conflict and cooperation between players in an improvisational scenario and the mathematical model we can use to frame it, he suggests the application of *game theory*.

Grounded in this theoretical framework for how such a music game might work, this project set out to explore a new method of algorithmic music composition. Thus, the goals of the project were:

- Implement in Haskell the game theoretic model outlined in Hudak’s paper

- Produce a meaningful, new way of algorithmically generating music that resembles human improvisation
- Allow for the implementation to be extended through user-defined ideas of payoffs, strategies and rules

2 Background: Game Theory

2.1 Concepts

Game theory is the study of strategic decision making, used to model interactions between agents whose decisions affect each other. In traditional game theory, a *game* is a situation in which two or more agents make decisions. The agents are called *players* and the decisions they make are *moves*. The *rules* of the game define a set of legal moves available to each player for any *state* of the game. The algorithms that players employ to choose moves at any game state are called *strategies* and can either be algebraically defined or non-deterministic. A player is said to be playing an optimal strategy when it always results in her receiving her maximum possible *payoff* at the game's conclusion, where the payoff is a quantified measure of the player's success. Game theory uses *game trees* to model the entire decision making process, where each node contains the state of the game, and each branch is an event that modifies that state.

2.2 Experimental Game Theory

Hagl is a domain-specific language embedded in Haskell that allows for simple and modular definitions of experimental games. It provides easy interfaces for defining what each of the above terms means in the context of a specific game. Each game is an instance of a type class, which requires the definition of a number of pieces.

```
class GameTree (TreeType g) => Game g where
  type TreeType g :: * -> * -> *
  type State g
  type Move g
  gameTree :: g -> (TreeType g) (State g) (Move g)
```

`TreeType`, `State`, and `Move` are associated types defined in terms of `g`, the particular game. A `Game` instance provides a `gameTree` function, which constructs the game tree for the specific game. Hagl provides a number of operations for interacting with the game values by examining internal tree nodes and edges.

The `TreeType` must fit into Hagl's representation of a `GameTree` as well. This game tree represents every possible sequence of moves from an initial game state to all final states. Game trees are rooted at the starting state, and each branch




Figure 1: Two levels of the Tic-tac-toe game tree

or edge represents a possible move for a player or chance for an external event to change the game state. Leaf nodes represent the final states, associated with the payoffs for every player. Game trees may either be *continuous* or *discrete*. In discrete game trees, there are a fixed number of edges and therefore a finite set of moves, while in continuous trees the set of moves is potentially infinite and is defined by a function from move to subtree.

When designing the `gameTree` function, one must consider the *information* each player has about the game state; this knowledge may or may not be shared amongst the players. When the players make their decisions at the same time, we call this a *simultaneous* game and restrict the players' knowledge of the others' decisions until all moves have been made and the game state updates en masse at the end of a round.

2.3 A formal treatment of Tic-tac-toe

As a brief example, consider the well-known game Tic-tac-toe. The players are the characters 'X', and 'O', who make their moves by placing their mark on a 3 by 3 grid. The rules state that the players must alternate turns. On each turn a player must put his mark in a previously unoccupied cell. When three cells in a row are occupied by the same player (horizontally, vertically, or diagonally), that player wins and the game is over. If the board fills up before this can happen, the result is a draw. Players each have knowledge of where both players have moved on the board prior to the current turn (therefore, Tic-tac-toe is not an example of a simultaneous game). Payoffs in this game are fixed: 1 for the winning player, -1 for the losing player, and 0 for each player in the case of a draw. The first few levels of the game tree might look like *Figure 1* above.

The first row models the options for player X's initial move. The second row shows O's possible follow-up moves.

3 Musical Improvisation as a Game

Given this general game theoretic framework, how do we then map the fundamental components of a game to aspects of an improvisational performance of two or more musicians?

3.1 Moves

For a musical game, moves take the form of fixed-duration, time-stamped musical events. Rather than thinking of a player’s full performance as a sequence of notes, we imagine it as a sequence of musical events, that is, a stream of his decisions of what and how to play at each time-stamp.

Events must necessarily be both time-stamped and of a fixed-duration to ensure that players make their decisions of what to play for each event at the same time. Notes in the musical sense are associated with a given duration, but consider a game in which one player begins playing a whole note and the other a quarter. When should we define the next decision point at which the players can choose their next moves? When the second player is ready to make another move, the first still has three more beats to play. Thus, we break notes into events of fixed-duration and allow a player to choose to extend the pitch of the previous event if he or she desires to play a note of a longer duration.

In his paper, Hudak notes that the formal model of musical interaction “is limited to controlling an instrument’s sound, for the purposes of realizing fundamental parameters such as pitch in addition to more subtle issues of articulation, dynamics and phrasing.” Indeed, in our simple implementation, we focus solely on controlling pitch, ignoring the more subtle parameters and keeping the players’ instruments fixed.

3.2 Payoff

The payoff of the game, used to measure a player’s success in the performance, translates to a player’s notion of musical aesthetic. This is a measure of how good the musician considers the sound of his own performance in the context of the whole peice. This preference is unique to each player and may be defined along any axis of the musical design space. The realm of musical aesthetic is practically infinite, and a full discussion of music theory of this depth is well beyond the scope of this paper.

3.3 Strategy

Each player in a musical game implements her own strategy, which determines what she will play at any given point in the performance. In a musical impro-

visational game, a strategy can be as simple as a player adhering to her given score, never improvising at all. Alternatively, a strategy can be more sophisticated. It may take into account past or anticipated payoffs, a player’s own past moves or those of another player, or some other musical strategy with the only restriction that any move generated fall within those allowed by the rules of the game.

For each individual, the goal of the game is to maximize his own payoff. However, we note that unlike zero-sum games, which are purely competitive, a musical game is more cooperative in nature. Indeed, there exists some competition between players – Hudak imagines two soloists “vying for attention.”. However, intuitively there is a point at which trying to make the other player “sound bad” (i.e. reduce his payoff) has an equally deleterious effect on one’s own performance. Therefore, the best outcome for each player individually is likely to be one in which the sum of their collective payoffs is also maximized.

3.4 Game Tree

We model this musical improvisational game with a discrete, simultaneous game tree. We eliminate the possibility of chance from the tree: at internal nodes, only a player’s choices may affect the state of game. Every edge represents a choice a performer can make to stick to or deviate from the score. The tree represents the mutual recursion between the players in that each player’s decision depends on the previous decisions of all other players.

4 Improvise: our implementation

In order to implement music as a playable game, we must implement the aforementioned components of games: moves, payoffs, and strategies. The choices made here define the framework for the game itself, and therefore, the range of music that can be produced from it. Our implementations for some components are static and fixed, and others can change with each instance of an `Improvise` game.

4.1 Static Components

4.1.1 Moves

The first important piece is the `MusicMv`:

```
data MusicMv = Begin Pitch
              | Extend Pitch
              | Rest
```

which defines a move as either a **Begin** of a **Pitch**, **Extend** of a **Pitch**, or a **Rest**. Each represents a musical event of the game's defined smallest duration. Here again we make the distinction between a **Begin** and **Extend**. A **Begin** represents the onset or attack of the given pitch, whereas an **Extend** is a continuation of the previous pitch without a distinct onset. The difference lies in the following example: Given a sequence of moves [**Begin** p, **Begin** p], the sound produced would be a note of length one unit followed by another note of length one. There will be an audible separation of the two notes and a distinct beginning of the second immediately following the termination of the first. In contrast, the move sequence [**Begin** p, **Extend** p], would produce a note of twice the duration with no audible onset beyond the first. The notes will sound as one continuous event lasting two units. A **Rest** represents no audible sound, simply one unit of silence.

The reasoning behind giving **Extend** a pitch is somewhat subtle. It is necessary in our implementation for an **Extend** to directly follow either a **Begin** or another **Extend**. In both cases, the extension of the preceding note must contain the same pitch. It is unclear then why **Extend** should be given a pitch at all, as it should be simple enough to look back through the previous moves until the most recent **Begin**, to which the extension is being applied, is found. We found this to be slightly impractical later on when attempting to write strategies and payoffs. The cost in storage of the **Pitch** must be weighed against the cost in time of repeatedly looking back to find it. This decision also lifts the burden from the programmer when writing payoffs and strategies.

4.1.2 State

The next piece of the game implementation is the **Performer**:

```
data Performer = Performer { realization :: [MusicMv]
                             , future      :: [MusicMv] }
```

represents an individual player. The **realization** is a list of musical events that have already occurred, i.e. the player's interpretation of the score thus far (most recent event first). The **future** is a list of the upcoming events, i.e. the remaining portion of her score. Therefore, the whole performance is a **Performer** list, given by:

```
type Performance = ByPlayer Performer
```

A **ByPlayer** a, as defined in Hagl, is a list of a which has length equal to the number of players in the game.

4.1.3 Game Tree

From here, we must design the generation of the tree to fit a few requirements. First, it must keep a state in each node, and continually modify that state as

moves get played. Second, it must appear to players that they are all making moves at once, despite the fact that a node can only represent a move by one player. Third, it must have leaf nodes populated with final payoff whenever the state makes it clear the game is over. Finally, it must be efficient and expand as needed rather than on first call of the function.

We use a recursive function to generate the tree and take advantage of lazy evaluation to prevent the whole game tree from being expanded at once. Nodes that are never explored in the process of the game will not be expanded at all. To enforce the stipulation in a simultaneous game that players be ignorant of each others' moves during a round, we build out the game tree by accumulating all of the moves for the round in a list and do not update the game state until every player has made her decision. This accumulating list of moves is never seen by a player's strategy, but only used in game tree generation, so is safe.

4.2 Dynamic Components

It is desirable to modularize those aspects of the game that we identify as particular to a unique game execution. Most obviously, we will want to experiment with different initial scores, that is, a user should be able to play an *Improvise* game with any song of his choosing.

We wrap these dynamic quantities in the type *Improvise*:

```
data Improvise =
  Imp { state      :: Performance
       , payoff    :: Performance -> Payoff
       , playable  :: Performance -> PlayerID -> [MusicMv]}
```

Here a user may define the number of players and their initial scores in *state*. In most games, prior to execution, each *Performer* will have a *future* loaded with an individual score as a list of *MusicMv* and an empty *realization*. In our implementation, we provide infrastructure for rendering a list of MIDI files of each player's score into an initial state of type *Performance*.

4.2.1 Payoff

Imp also requires a function *payoff* that generates a Hagl *Payoff* matrix for any given game state. *payoff* is a user-defined function that generates a Hagl *Payoff* matrix for any given game state. At various times in the development of *Improvise*, it was suggested that *payoff* be generated based on tempo changes, sequences of notes, or even by more sophisticated schemes derived from common jazz improvisational techniques. Recognizing the fact that there are practically an infinite number of dimensions upon which to judge musical aesthetic, it is critical that *payoff* generation has the most general type possible.

In our implementation, we modeled an extremely simple idea of how a player might judge the sound of his performance: that of pitch intervals, the relative distance on the scale between two notes. In Western music theory, some pitch intervals are generally considered to sound “consonant”, or pleasing - the classic example is that of a major-third, two notes that are four semitones apart on a scale. Other intervals, like the augmented-fourth (six semitones apart), may be termed “dissonant” and are considered to cause tension in a piece. Most pieces will use a mix of consonant and dissonant sounds to alternately build and resolve tension.

For an **Improvise** game using the **intervalPayoff** function, this means that each player provides his own interval preferences: an association list of intervals (as a integer number of half-steps, or semitones) and a **Float**, representing the relative value the player places on playing that interval. Positive payoffs denote favorable intervals for the player, while negative payoffs signify undesirable ones.

Intervals not included in the list have a baseline value of zero for the player; an empty interval preference list denotes a player who values all intervals equally. We also distinguish the “top” player, playing the higher note of the interval from the “bottom” player playing the low note by allowing both positive and negative intervals. For example a player with the preferences [(4, 4.0), (-4, 2.0)] always values major-thirds in a piece, but prefers to be the bottom player. The idea is that if two players both have positive values for the same interval, they will collude to play that interval more often. Alternatively, a mix of positive and negative payoffs for a given interval results in a more competitive relationship between players.

4.2.2 Rules

Finally, **playable** is a function for generating the legal moves for a given player for any state of the game. For our studies, we defined a simple **playable** function derived from the player’s score: legal moves are those within a set number of semitones away from the pitch given in the score. We called this a “range-limited” move generation scheme. To make the scheme slightly more sophisticated, we also allowed players to “look back” at their most recently played note and play pitches within a range around it as well. Players always have the option of resting (emitting no sound) but we also stipulate that an **Extend** move is only legal following a previous **Begin** or a previous **Extend**, never after a **Rest**.

This **playable** function is fairly naive in the realm of music improvisation and many other more realistic schemes have been suggested. Given the generality of the **playable** function in the implementation of **Improvise**, it should be possible to generate moves based on other musical attributes of the piece.

4.2.3 Strategies

Hagl comes prepackaged with a classic game theory strategy called `minimax`, in which a player simultaneously endeavors to maximize his own payoff and diminish that of his adversary. To more closely model the cooperative nature of a musical game, we implemented a variation on this strategy, which we called `maximize`. In this strategy, the player whose turn it is calculates an intermediate payoff for each of his possible moves. From these, we pick the three nodes that yield the best intermediate payoff, and explore those nodes by recursively calling the same `maximize` strategy for the next player. The recursive algorithm continues until the bottom of the tree or a defined depth limit, at which point, the player who began the exploration considers the three payoffs at every level, and choose the maximum from these.

The reason it is important to explore the tree rather than just picking the move that yields the highest intermediate payoff is that another player's decision in the future may affect your final payoff, and making an optimal move now may cause a player to fall into a bad situation in the future. The payoff is intermediate because it is not necessarily indicative of the final payoff.

4.3 The State Space

The design decisions made throughout the implementations of the above dynamic components are centered around the state space. The size of a tree grows exponentially with its branching factor. This means that if each player makes n decisions each from m choices, there are on the order of m^n nodes in the tree. The decision to limit the range of possible pitch deviation from the score in the `limitByRange` function aims to limit m , while the depth limit `maximize` aims to limit n . If one was to make every note on the piano legal in every context, the strategies may explore 60 million game states within four decisions. Our strategies take advantage of the lazy evaluation of the game tree, and limit the branching factor by limiting possible moves from each state. All further implementations of strategies and `playable` should keep these factors in mind, though we have kept these as dynamic parameters as the willingness to allocate computational resources in the pursuit of good improvisation may vary from person to person.

5 Case Studies

In this section we will walk through examples that illustrate some of the capabilities of the `Improvise` game. The dynamic components used in the examples are fairly straightforward and should serve to demonstrate what can be accomplished with a basic strategy, payoff, and move domain. For these examples, we will often refer to a payoff function `pay` defined as follows:

```

prefs1 = [(-3, 2), (-5, 2), (5, 2), (3, 2)]
prefs2 = [(5, 1), (3, 1)]
pay = intervalPayoff [prefs1, prefs2]

```

As a reminder, this uses the association lists of intervals and preference weights to generate a payoff matrix for each final state.

The two strategies showcased are `maximize pay`, which uses our maximize strategy given the payoff function, and `justTheScore`, which as its name indicates, plays the original score.

The `playImprovise` function used to execute these games takes all the dynamic pieces we previously discussed:

```

playImprovise ::
  (Performance -> Payoff) --payoff function
-> Performance           --starting score
-> (Performance
    -> PlayerID
    -> [MusicMv])         --playable moves
-> [Hagl.Player Improvise] --players in the game (strategies)
-> IO ()

```

5.1 “Mary Had a Little Lamb”

In this first example we use “Mary Had A Little Lamb” as our starting score for both players.

```

playImprovise (intervalPayoff [prefs1, prefs2])
  ByPlayer [mary, mary]
  (limitByRange 2)
  [justTheScore, (maximize pay)]

```

Below is the musical output of this example on the left as a list of moves, annotated with the intervals and payoffs for each player on the right.

["Mr. Score ", "Missus Maxine "]	Interval1	Payoff1	Interval2	Payoff2
[Begin (A,4), Begin (A,4)]	0		0	
[Extend (A,4), Extend (A,4)]	0		0	
[Begin (G,4), Begin (G,4)]	0		0	
[Extend (G,4), Extend (G,4)]	0		0	
[Begin (F,4), Begin (F,4)]	0		0	
[Extend (F,4), Extend (F,4)]	0		0	
[Begin (G,4), Begin (E,4)]	-3	2	3	1
[Extend (G,4), Extend (E,4)]	-3	2	3	1
[Begin (A,4), Begin (Fs,4)]	-3	2	3	1
[Extend (A,4), Extend (Fs,4)]	-3	2	3	1
[Begin (A,4), Begin (Fs,4)]	-3	2	3	1

[Extend (A,4), Extend (Fs,4)]	-3	2	3	1
[Begin (A,4), Begin (Fs,4)]	-3	2	3	1
[Extend (A,4), Extend (Fs,4)]	-3	2	3	1
[Extend (A,4), Extend (Fs,4)]	-3	2	3	1
[Extend (A,4), Extend (Fs,4)]	-3	2	3	1

Payoffs:
[20, 10]

Going through this output, we see that at first both players choose to play in unison, both following the score they were given. This is due to the limited range by which they are allowed to deviate from the score. To this point, they have not had an opportunity to produce any intervals of any worth in terms of payoff. However, when we reach the 7th move, we see player 2 play an E below player 1's G. The interval from player 1's perspective is -3 as the move from a G to an E is down 3 half-steps. Conversely, the interval from player 2's perspective is an E to a G which is an upward distance of 3 half-steps. Player 1 values the interval -3 at a payoff of 2, while player 2 values their interval, 3, at 1.

We can see that for all of the following moves, the intervals produced are exactly the same. This is quite logical given the input to the game. Looking back at **prefs1** and **prefs2**, we can see that player 1 and player 2 have both achieved the maximum payoff possible. However, both players give equal weight to a 5 (or -5) interval as +/-3, so why would player2 not choose to produce any? This turns out to be due to the tie-breaking in the **maximize** strategy. We chose to have ties in payoff broken by playing the note closest to the original score. This prevents players from deviating further and further on each successive move, creating notes in entirely different registers that do not remotely resemble the score. Consequently, because player 1 plays only the score, intervals of 3 half-steps are by definition closer to the score than those of 5 half-steps.

Though the choices made in terms of payoffs for this example are rather simple and not necessarily the most meaningful for producing complex or interesting music, they serve to illustrate the functionality of the preference-based payoff system and the effectiveness of the **maximize** strategy.

The example preferences used could be expanded to include more intervals, or even the whole range of -12 to +12. Additionally, a player can assign a relative preference between different intervals by assigning them different weights. For example, if we doubled the weight of 5 and -5 in the above preferences, it is likely we would have seen more of them in the output. Though ties are broken by playing closest to the score, these payoffs would not create a tie, there would be a clearly superior move available according to the payoffs.

5.2 “Don’t Stop Believing”

While Mary Had a Little Lamb produces interesting and pleasing results, it is not often that two players start out playing and improvising from the same score. What happens in a more realistic scenario, where they’re already playing a pleasing harmony? We tested this with the first few measures of “Don’t Stop Believing” by Journey.

```
playImprovise pay
  ByPlayer [DontStopMiddle, DontStopBass]
    (limitByRange 2)
    [justTheScore, maximize pay]
```

Once again, we have player 1 adhering to the score while player 2 improvises, using the same payoff function and set of preferences as above. However, because these players are already harmonizing, even without any improvisation, they already each achieve a positive payoff of [16,4]. However, with improvisation, they have a payoff of [8,15], and the realization sounds similar, but with a slightly different bass line. Player 2 has kept in mind a goal of making the whole performance better – the cumulative payoff is 23 rather than 20 – while still working to focus on improving the score in terms of his preferences.

6 Conclusion

The field of algorithmic music generation is a relatively new and growing one. A multitude of tools and methods exist, each employing a different model for music and its creation. While many of these tools are quite sophisticated and capable, none appear to implement a game theoretic model. The model used in **Improvise** represents an approach that more closely models the way humans make decisions and interact in music. It also offers opportunity for near-infinite customization and extension. Customized payoffs and strategies can be easily integrated to produce music according to any definition of musical aesthetic and run on any score. Because this work is in a largely unexplored area, there is much room for expansion and alternative representations.

It is natural to model a music game discretely, as there are typically a limited number of moves a person would consider in their search. In a continuous game tree, you can ask the tree whether a move from a given node is acceptable, but you must have a sense of the moves of which you ask the tree about rather than getting the list from the edge set that a discrete tree node contains. This may prove an interesting way to represent the tree, as players do typically have a sense of what notes would sound good, and do not need to explore the tree in order to find them. This could heavily limit the explored state space according to how adventurous the strategies ended up being. This would involve a nearly complete reimplementaion of this work, as the game tree and representation of

legal moves are what provide the structure for the current formulation of the game.

There is also much work to be done in defining how payoff calculations should happen. Calculating a payoff based on preferences of intervals is simple, and it is already widely acknowledge that there are intervals that sound more pleasing than others. However, the aesthetic of entire pieces of music is much more based upon ideas of repetition and variation or rates of consonance to dissonance than only harmonies. These types of calculations are possible, but involve a much more specific identification of aesthetic of music, which cannot be universally defined. This project did not aim to explore those areas, but this may be an area for future exploration by those with technical backgrounds in music theory.

There is also space to explore strategies, which are easier to expand upon and reason about computationally. A strategy is a state monad transformer, so it can hold on to and maintain awareness of the repetition and variation it wishes to accomplish, if one were to define an efficient data structure. We can be more confident that this is computationally feasible as long as exploration of the exponentially expanding game tree is kept limited according to predefined desires, but this is clearly also a place where much work would be needed to make progress.

7 Acknowledgements

Paul Hudak

Eric Walkingshaw

Norman Ramsey

Advanced Functional Programming Classmates